

Predicting Customer Churn for SyriaTel Telecommunications Company

Student names:

1. Peggy Obam
2. Rael Ndonye
3. Alan Omondi
4. Jackson Maina
5. Jared Bii
6. Brenda Kinoti

Student pace: part time

Introduction:

The goal of this project is to develop a predictive model that can accurately forecast customer churn for SyriaTel, a telecommunications company. By identifying customers who are likely to churn, SyriaTel can take proactive measures to retain them, thereby reducing revenue loss and improving customer satisfaction. This proposal outlines the project's objectives, dataset selection, methodology, and expected deliverables.

Objectives:

The primary objectives of this project are as follows:

- a. To Build a classification model to predict customer churn for SyriaTel.
- b. To Identify the key factors influencing customer churn.
- c. To Provide insights and recommendations to SyriaTel for effective churn management.

Dataset Selection:

For this project, we have chosen the "SyriaTel Customer Churn" dataset. The dataset provides a comprehensive set of customer-related features that can be used to analyze and predict customer churn. The features of this dataset provide valuable insights into customer behavior, usage patterns, and account details. By analyzing this data, we aim to develop a predictive model that can identify customers who are likely to churn. By leveraging the available features, such as call duration, usage patterns, and customer service interactions, we can gain a better understanding of the factors contributing to customer churn and explore potential strategies to reduce churn rates. This dataset is particularly suitable for our objectives, as it provides the necessary information to understand customer behavior and predict churn.

Methodology:

The project will follow the following steps:

- a. Exploratory Data Analysis: We will perform an in-depth exploration of the dataset to gain insights into the distribution of variables, identify patterns, and detect any data quality issues.

- b. Data Preprocessing: This step involves handling missing values, encoding categorical variables, and scaling numerical features. We will also address any outliers or data inconsistencies to ensure the reliability of our analysis.
- c. Feature Selection: We will identify relevant features that have a significant impact on customer churn prediction.
- d. Model Selection and Training: We will compare various classification algorithms, such as logistic regression, decision trees, and random forests, to select the most suitable model for predicting customer churn. The chosen model will be trained using the labeled dataset.
- e. Model Evaluation: We will assess the performance of the trained model using appropriate evaluation metrics, including accuracy, precision, recall, and F1-score. This step will help us understand how well the model predicts customer churn.
- f. Model Optimization: We will fine-tune the selected model by adjusting hyperparameters and employing techniques like grid search. This optimization process aims to maximize the model's predictive capabilities.
- g. Interpretation and Reporting: The results of our analysis will be interpreted and presented in a comprehensive report. We will provide actionable insights and recommendations for the company to implement.

A.Data Exploration

In [34]:

```
#importation of necessary libraries and loading of the data set.
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import math
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
import warnings
warnings.filterwarnings("ignore")
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, conf
from sklearn.metrics import f1_score
from imblearn.over_sampling import SMOTE, ADASYN
df=pd.read_csv('data.csv')
```

In [35]:

▶

```
df.head() #View the first few rows of the dataset
```

Out[35]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...

5 rows × 21 columns

In [36]:

▶

```
# Checking the dimensions of the dataset  
print("Shape of the dataset:", df.shape)
```

Shape of the dataset: (3333, 21)

In [37]:

▶

```
# Checking the column names  
print("Column names:", df.columns)
```

Column names: Index(['state', 'account length', 'area code', 'phone number',
'international plan', 'voice mail plan', 'number vmail messages',
'total day minutes', 'total day calls', 'total day charge',
'total eve minutes', 'total eve calls', 'total eve charge',
'total night minutes', 'total night calls', 'total night charge',
'total intl minutes', 'total intl calls', 'total intl charge',
'customer service calls', 'churn'],
dtype='object')

Observations:

The column names include various customer-related information such as 'state', 'account length', 'area code', 'phone number', 'international plan', 'voice mail plan', 'number vmail messages', and several other features related to call duration, charges, and customer service interactions. This suggests that the dataset covers a wide range of customer attributes.

In [38]:



```
# Getting summary statistics of numerical features  
print(df.describe())
```


	account length	area code	number vmail messages	total day minut
es \				
count	3333.000000	3333.000000	3333.000000	3333.0000
00				
mean	101.064806	437.182418	8.099010	179.7750
98				
std	39.822106	42.371290	13.688365	54.4673
89				
min	1.000000	408.000000	0.000000	0.0000
00				

Observations:

From the summary statistics above, we can see that:
The average account length is approximately 101, with a minimum of 1 and a maximum of 243. The average total day minutes is around 179.8, with a standard deviation of 54.5. The average total eve minutes is approximately 201.0, with a standard deviation of 50.7. The average total intl calls is about 4.5, with a maximum of 20. The average total intl charge is about 11.6, with a maximum of 30.91.

In [39]:

total day calls	total day charge	total eve minutes	total eve cal
-----------------	------------------	-------------------	---------------

ls \				
# Checking the data types of columns				
print(df.dtypes)	3333.000000	3333.000000	3333.000000	3333.0000
00				
mean	100.435644	30.562307	200.980348	100.1143
state	object			
account length	100.069084	int64	50.713844	19.9226
area code	int64			
phone number	0.000000	object	0.000000	0.0000
international plan	object			
voice mail plan	0.000000	object	166.600000	87.0000
number vmail messages	int64			
total day minutes	float64	50.500000	201.400000	100.0000
total day calls	int64			
total day charge	float64	56.790000	235.300000	114.0000
total eve minutes	float64			
total eve calls	int64	59.640000	363.700000	170.0000
total eve charge	float64			
total night minutes	float64			
total night calls	int64			
total night charge	float64	3333.000000	3333.000000	
total intl minutes	float64	200.872037	100.107711	
total intl calls	int64	50.573847	19.568609	
total intl charge	float64	23.200000	33.000000	
customer service calls	int64	167.000000	87.000000	
churn	bool	201.200000	100.000000	
type: object	20.000000	235.300000	113.000000	
max	30.910000	395.000000	175.000000	

Observations:

The data types of the columns indicate that most features are represented as integers or floats, while a few are categorical variables (object), such as 'state', 'international plan', and 'voice mail plan'. The 'churn' column is a boolean variable representing whether a customer has churned (True) or not (False).

total intl charge	customer service calls
count	3333.000000
mean	2.764581

```
std          0.753773          1.315491
In[40]:      0.000000          0.000000
min          2.300000          1.000000
25%          2.780000          1.000000
# Checking for missing values
50%          3.270000          2.000000
print("Missing values:\n", df.isnull().sum())
75%          5.400000          9.000000
max          5.400000          9.000000
```

Missing values:

state	0
account length	0
area code	0
phone number	0
international plan	0
voice mail plan	0
number vmail messages	0
total day minutes	0
total day calls	0
total day charge	0
total eve minutes	0
total eve calls	0
total eve charge	0
total night minutes	0
total night calls	0
total night charge	0
total intl minutes	0
total intl calls	0
total intl charge	0
customer service calls	0
churn	0

dtype: int64

Observations:

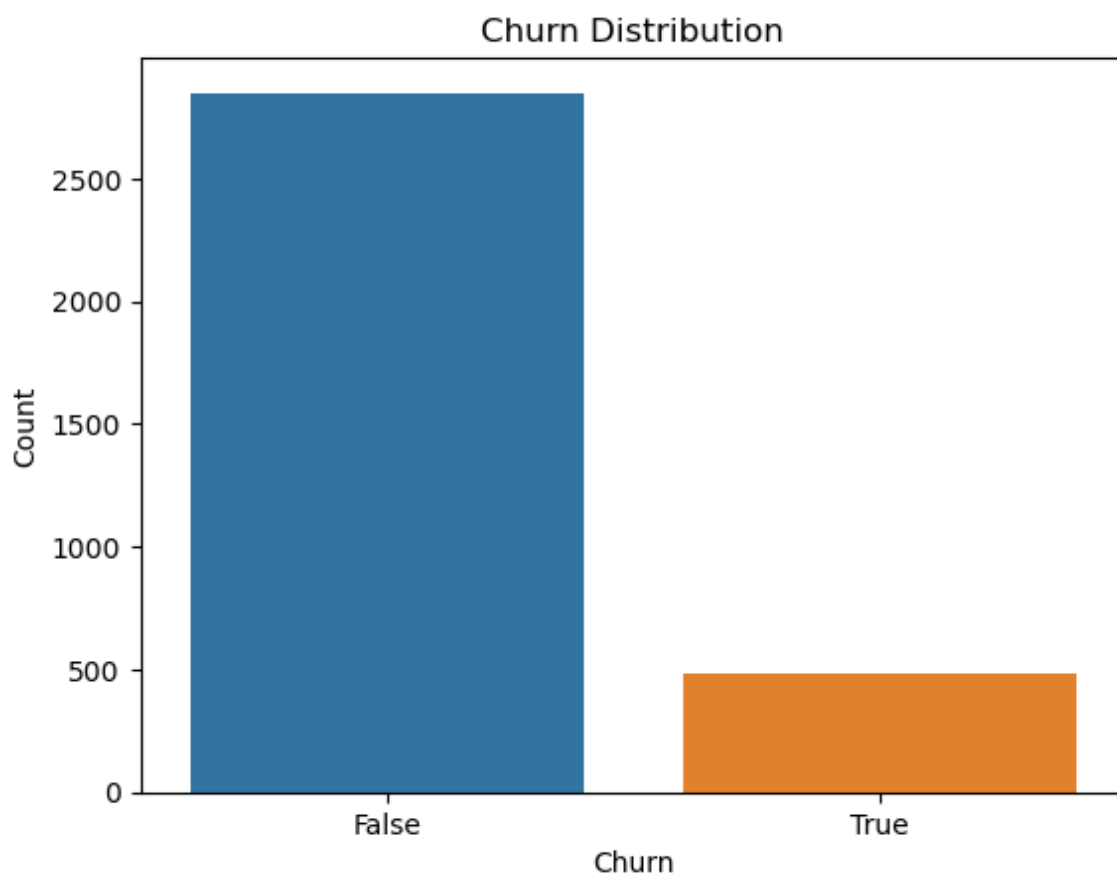
As seen above,the dataset contains no missing values.This suggests that the dataset is complete, with no null or missing entries in any of the columns.This is advantageous as it ensures that the data is ready for further analysis and modell ing without the need for imputation or handling missing data. It provides a reliable foundation for exploring relationships between variables and deriving meaningful insights from the data.

Explolatory Data Analysis

Univariate analysis:

In [41]:

```
#Histogram of Churn Distribution:  
#To visualize the distribution of the target variable 'churn':  
sns.countplot(x='churn', data=df)  
plt.xlabel('Churn')  
plt.ylabel('Count')  
plt.title('Churn Distribution')  
plt.show()
```



Observations:

-The majority of customers in the dataset did not churn (represented by 'False' in the 'churn' variable). -The number of churned customers is noticeably smaller compared to the number of customers who did not churn.

In [42]:



```
#Box Plot of Customer Service Calls:  
#To analyze the distribution and potential outliers of the 'customer service calls' variable  
plt.boxplot(df['customer service calls'])  
plt.xlabel('Customer Service Calls')  
plt.ylabel('Count')  
plt.title('Box Plot of Customer Service Calls')  
plt.show()
```



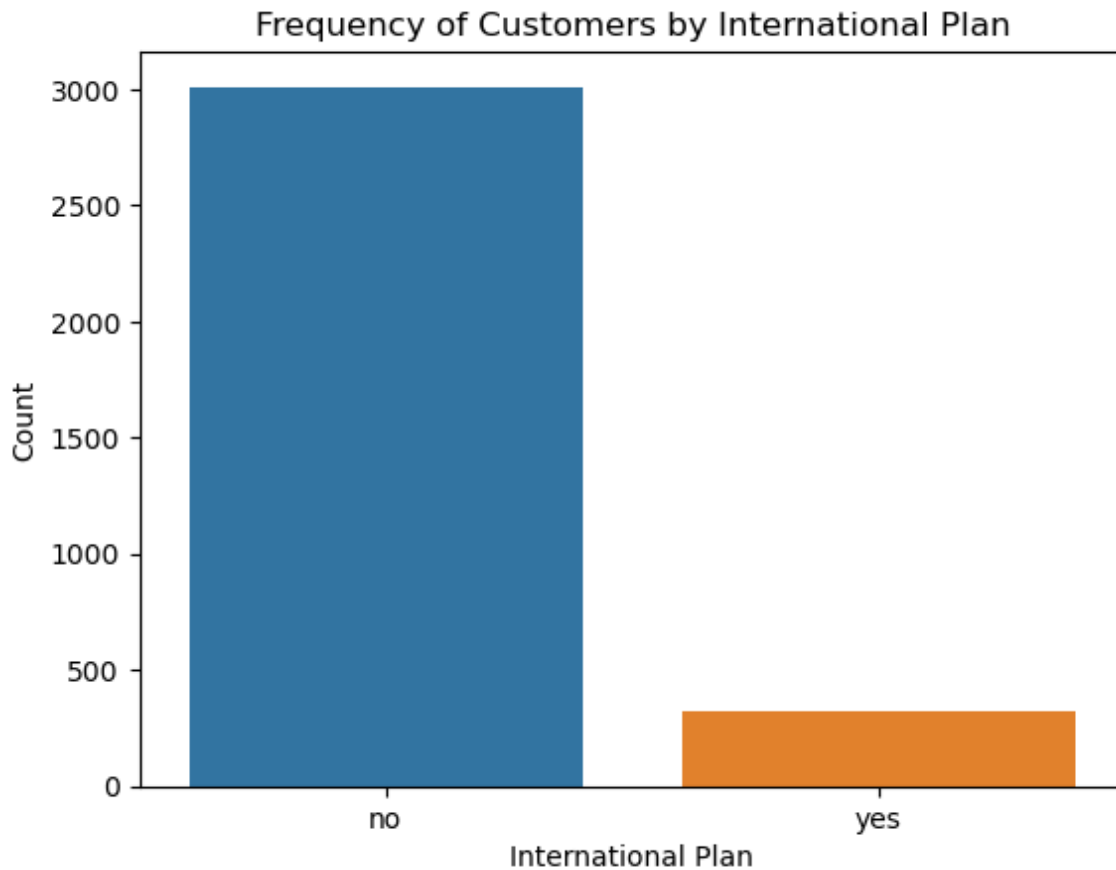
Observations:

-The box plot shows the distribution of the 'customer service calls' variable. -The median number of customer service calls is around 1.5. -There are a few outliers with a higher number of customer service calls.

In [43]:



```
#Bar Plot of International Plan:  
#To examine the frequency of customers with and without the 'international plan':  
sns.countplot(x='international plan', data=df)  
plt.xlabel('International Plan')  
plt.ylabel('Count')  
plt.title('Frequency of Customers by International Plan')  
plt.show()
```



Observations:

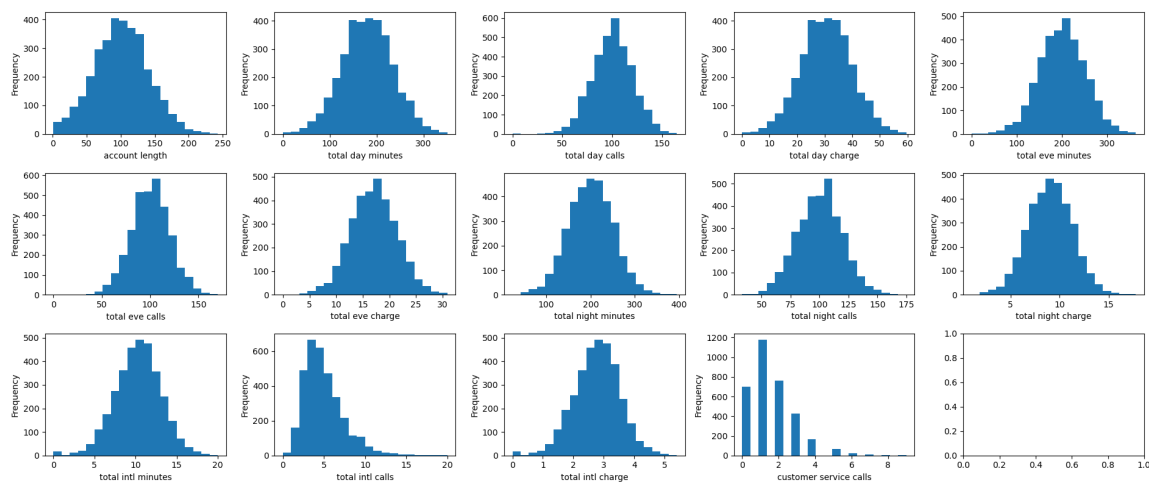
-The bar plot displays the frequency of customers with and without an international plan. -The majority of customers do not have an international plan. -A smaller proportion of customers have opted for an international plan.

In [44]:

```
# Distribution Plots for all features
fig, axes = plt.subplots(nrows=3, ncols=5, figsize=(19, 8))
variables = ['account length', 'total day minutes', 'total day calls', 'total day charge',
            'total eve minutes', 'total eve calls', 'total eve charge', 'total night minutes',
            'total night calls', 'total night charge', 'total intl minutes', 'total intl calls',
            'total intl charge', 'customer service calls']

for i, variable in enumerate(variables):
    ax = axes[i // 5, i % 5]
    ax.hist(df[variable], bins=20)
    ax.set_xlabel(variable)
    ax.set_ylabel('Frequency')

plt.tight_layout()
plt.show()
```



Observations

Most of the features are normally distributed except for **total intl calls** and **customer service calls** which are left skewed.

Bivariate Analysis:

In [45]:

```
# Compute the correlation matrix
corr = df.corr().abs()

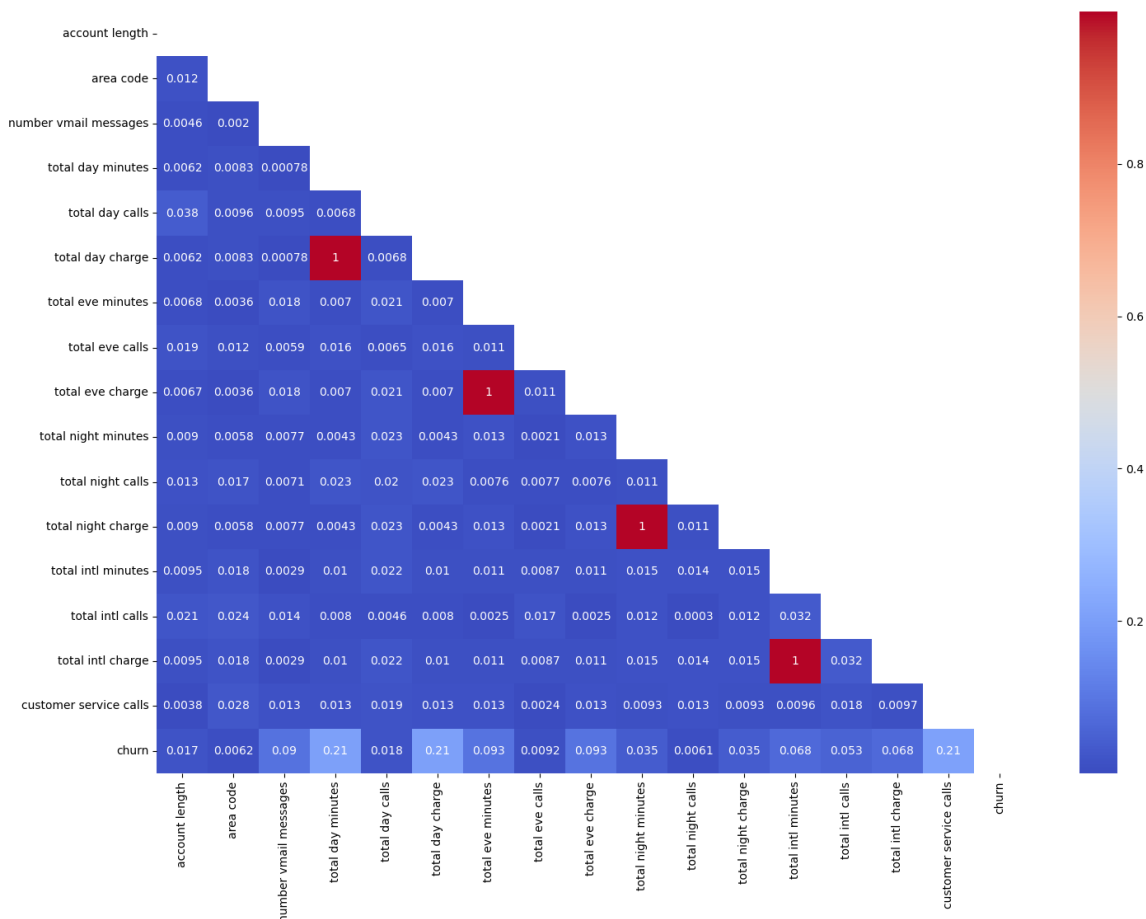
# Generate a mask for the upper triangle
mask = np.triu(np.ones_like(corr, dtype=bool))

# Set up the figure and axes
fig, ax = plt.subplots(figsize=(17, 12))
fig.suptitle('Correlation Matrix of SyriaTel Customer Churn data', fontsize=30, y=0.95, f

# Create a heatmap using the mask for the upper triangle
heatmap = sns.heatmap(corr, mask=mask, cmap='coolwarm', annot=True)

# Show the heatmap
plt.show()
```

Correlation Matrix of SyriaTel Customer Churn data



- The correlation between total day minutes and total day charge is 1, suggesting a strong linear relationship. It is possible to remove one of these variables without significant loss of information.
- Similarly, the correlation between total eve minutes and total eve charge is 1, indicating a strong linear relationship. One of these variables can be dropped.
- The correlation of 1 between total night minutes and total night charge implies a strong linear relationship, allowing us to drop one of these variables.
- Likewise, the correlation between total intl minutes and total intl charge is 1, indicating a strong linear relationship, allowing for the removal of one of these variables.

Due to high multicollinearity between each other, we dropped the following columns:

- total day minutes
- total eve minutes
- total night minutes
- total intl minutes

In [46]:

```
remove = ["total day minutes", "total eve minutes", "total night minutes", "total intl minutes"]
df = df.drop(remove, axis = 1)
```

In [47]:

```
df.dtypes
```

Out[47]:

```
state                object
account length      int64
area code           int64
international plan   object
voice mail plan      object
number vmail messages  int64
total day calls      int64
total day charge     float64
total eve calls      int64
total eve charge     float64
total night calls    int64
total night charge   float64
total intl calls     int64
total intl charge    float64
customer service calls int64
churn                bool
dtype: object
```

In [48]:

```
# creating dummy variables
df2 = pd.get_dummies(df, columns=['state', 'international plan', 'voice mail plan'], drop_first=True)
```

In [49]:

```
df2["churn"] = df2["churn"].map({True: 1, False: 0})
```

In [50]:



```
# Print the updated DataFrame  
df2.info()
```


Data columns (total 65 columns):

#	Column	Non-Null Count	Dtype
0	account length	3333 non-null	int64
1	area code	3333 non-null	int64
2	number vmail messages	3333 non-null	int64
3	total day calls	3333 non-null	int64
4	total day charge	3333 non-null	float64
5	total eve calls	3333 non-null	int64
6	total eve charge	3333 non-null	float64
7	total night calls	3333 non-null	int64
8	total night charge	3333 non-null	float64
9	total intl calls	3333 non-null	int64

BASELINE MODEL- LOGISTIC REGRESSION MODEL

Logistic regression is a regression analysis technique that is specifically designed for situations where the dependent variable is categorical and can only take discrete values. It is used to estimate the probability of a particular event occurring. In this type of regression, we first identify our target variable (Y) and the predictor variables (X) that we want to analyze and understand their relationship with the target variable.

15	state_AZ				3333	non-null	uint8
16	state_CA				3333	non-null	uint8
17	state_CO				3333	non-null	uint8
18	state_CT				3333	non-null	uint8
19	state_DC				3333	non-null	uint8
20	state_DE				3333	non-null	uint8
21	state_FL				3333	non-null	uint8
22	state_GA				3333	non-null	uint8
23	state_HI				3333	non-null	uint8
24	state_IA				3333	non-null	uint8
25	state_ID				3333	non-null	uint8
26	state_IL				3333	non-null	uint8
27	state_IN				3333	non-null	uint8
28	state_KS				3333	non-null	uint8
29	state_KY				3333	non-null	uint8
30	state_LA				3333	non-null	uint8
31	state_MA				3333	non-null	uint8
32	state_MD				3333	non-null	uint8
33	state_ME				3333	non-null	uint8
34	state_MI				3333	non-null	uint8
35	state_MN				3333	non-null	uint8
36	state_MO				3333	non-null	uint8
37	state_MS				3333	non-null	uint8
38	state_MT	0.524795	0.068627	0.490196	0.666667	0.755704	0.562359
39	state_NC	0.438017	0.068627	0.509804	0.745455	0.460597	0.605882
40	state_ND	0.561983	0.068627	0.000000	0.690909	0.693836	0.647059
41	state_NE	0.342975	0.000000	0.000000	0.490909	0.853454	0.517647
42	state_NH	0.305785	0.068627	0.000000	0.684848	0.475184	0.711764
43	state_NJ				3333	non-null	uint8
44	state_NM				3333	non-null	uint8
45	state_NV				3333	non-null	uint8
46	state_NY				3333	non-null	uint8
47	state_OH				3333	non-null	uint8
48	state_OK				3333	non-null	uint8
49	state_OR				3333	non-null	uint8
50	state_PA				3333	non-null	uint8
51	state_RI				3333	non-null	uint8
52	state_SC				3333	non-null	uint8
53	state_SD				3333	non-null	uint8
54	state_TN				3333	non-null	uint8
55	state_TX				3333	non-null	uint8


```

56 state_UT 3333 non-null uint8
In[53]: 57 state_VA 3333 non-null uint8
58 state_VT 3333 non-null uint8
#perform train test split
59 state_WA 3333 non-null uint8
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
60 state_WI 3333 non-null uint8
61 state_WV 3333 non-null uint8
# Create a logistic regression model using scikit Learn
62 state_WY 3333 non-null uint8
logreg = LogisticRegression(fit_intercept=False, C=1e12, solver='liblinear')
63 international_plan_yes 3333 non-null uint8
64 voice_mail_plan_yes 3333 non-null uint8
# fit the model on the training data
dtypes: float64(4), int64(9), uint8(52)
logreg.fit(X_train, y_train)
memory usage: 507.9 KB

```

Out[53]:

▼	LogisticRegression
LogisticRegression(C=1000000000000.0, fit_intercept=False, solver='liblinear')	

In [54]:

```

# Generate predictions
y_hat_train = logreg.predict(X_train)
y_hat_test = logreg.predict(X_test)

```




In [55]:

```
def evaluate_logistic_regression(X, y, test_size=0.2, random_state=42):
    # Split the data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.20, random_state=random_state)

    # Create a logistic regression model
    logreg = LogisticRegression(fit_intercept=False, C=1e12, solver='liblinear')

    # Train the model on the training data
    logreg.fit(X_train, y_train)

    # Predict the target variable on the test data
    y_hat_test = logreg.predict(X_test)

    # Calculate accuracy
    accuracy = accuracy_score(y_test, y_hat_test)

    # Calculate precision
    precision = precision_score(y_test, y_hat_test)

    # Calculate recall
    recall = recall_score(y_test, y_hat_test)

    # Calculate F1 score
    f1 = f1_score(y_test, y_hat_test)

    # Create a confusion matrix
    cm = confusion_matrix(y_test, y_hat_test)
    # Calculate the ROC curve
    y_hat_prob = logreg.predict_proba(X_test)[:, 1]
    fpr, tpr, thresholds = roc_curve(y_test, y_hat_prob)

    # Calculate the area under the ROC curve
    auc = roc_auc_score(y_test, y_hat_prob)

    # Plot the ROC curve
    plt.figure(figsize=(25,8))
    plt.subplot(1,2,1)
    plt.plot(fpr, tpr)
    plt.xlabel("False Positive Rate")
    plt.ylabel("True Positive Rate")
    plt.title("ROC Curve")

    # Create a heatmap of the confusion matrix
    plt.subplot(1,2,2)
    cm = confusion_matrix(y_test, y_hat_test)
    sns.heatmap(cm, annot=True, cmap='Blues', fmt='d')

    # Add labels and title to the plot
    plt.xlabel('Predicted')
    plt.ylabel('Actual')
    plt.title('Confusion Matrix')

    # Display the plot
    plt.show()
    # Assuming the evaluation dictionary contains the evaluation metrics
    evaluation = {
        "Accuracy": accuracy,
        "Precision": precision,
```

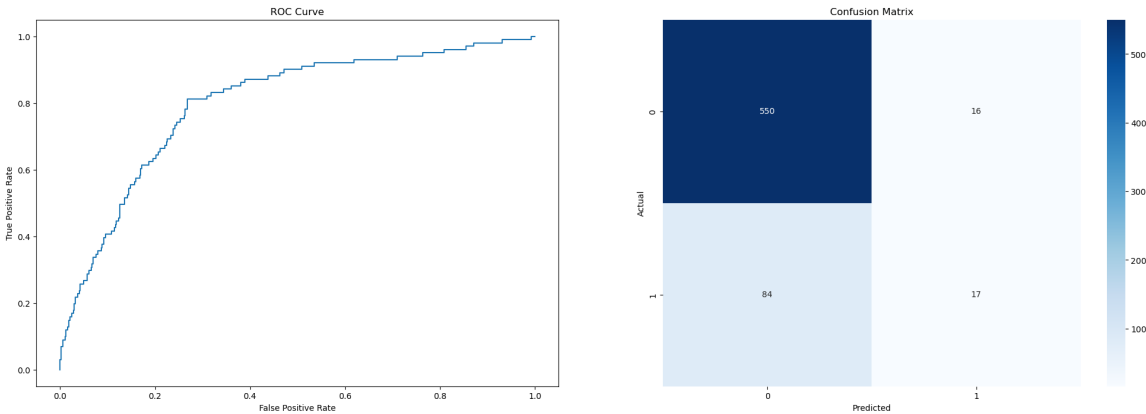
```
"Recall": recall,
"F1 Score": f1,
"Confusion Matrix": cm,
"AUC Score": auc
}

# Create a DataFrame from the evaluation dictionary
evaluation_df = pd.DataFrame.from_dict(evaluation, orient='index', columns=['Metrics'])

# Display the DataFrame
print(evaluation_df)
```

In [56]:

```
# Assuming you have X and y defined
evaluation = evaluate_logistic_regression(X, y)
evaluation
```



	Metrics
Accuracy	0.850075
Precision	0.515152
Recall	0.168317
F1 Score	0.253731
Confusion Matrix	[[550, 16], [84, 17]]
AUC Score	0.799181

Observations

1. Accuracy: The model's accuracy is 85%, indicating the percentage of correctly predicted instances. However, it should be considered alongside other metrics for a complete evaluation.
2. Precision: The precision is 51.5%, implying that only half of the predicted positive instances are actually true positives.
3. Recall: The recall is 16.83%, indicating the model's ability to correctly identify positive cases among all actual positive cases.
4. F1 Score: The F1 score, at 0.253731, represents a moderate balance between precision and recall.
5. Confusion Matrix [[550, 16], [84, 17]]: The confusion matrix provides a detailed breakdown of the model's predictions, including the number of true negatives, false positives, false negatives, and true positives.
6. AUC Score: The AUC score of 0.799181 suggests reasonable discrimination ability in distinguishing between positive and negative instances.

Overall, the observations reveal limitations in correctly identifying positive instances (low recall) and achieving a balanced precision and recall (low F1 score). Further analysis and model refinement may be

The confusion matrix provides information on the model's performance in classifying instances. In this case, the confusion matrix reveals the following:

- True Negatives (TN): The model correctly predicted "not churn" (0) for 550 instances where the actual value is also "not churn" (0).
- False Positives (FP): The model incorrectly predicted "churn" (1) for 16 instances where the actual value is "not churn" (0).
- False Negatives (FN): The model incorrectly predicted "not churn" (0) for 84 instances where the actual value is "churn" (1).
- True Positives (TP): The model correctly predicted "churn" (1) for 17 instances where the actual value is also "churn" (1).

Based on the confusion matrix, it is evident that the model struggles with correctly identifying instances that are actually churned, as indicated by the relatively high number of false negatives (84). This observation aligns with the low recall score (16.83%) obtained in the evaluation results.

Dealing with class imbalance: SMOTE

In [60]:



```
# # Previous original class distribution
# print(y_train.value_counts())

# # Fit SMOTE to training data
# X_train_resampled, y_train_resampled = SMOTE().fit_resample(X_train, y_train)

# # Preview synthetic sample class distribution
# print('\n')
# print(pd.Series(y_train_resampled).value_counts())
```

```
0    2284
1     382
Name: churn, dtype: int64
```

```
0    2284
1    2284
Name: churn, dtype: int64
```




In [61]:

```
# def evaluate_logistic_regression(X, y, test_size=0.2, random_state=42):
#     # Split the data into training and testing sets
#     X_train_resampled, X_test, y_train_resampled, y_test = train_test_split(
#         X, y, test_size=0.20, random_state=random_state)

#     # Create a logistic regression model
#     logreg_resampled = LogisticRegression(fit_intercept=False, C=1e12, solver='liblinear')

#     # Train the model on the training data
#     logreg_resampled.fit(X_train_resampled, y_train_resampled)

#     # Predict the target variable on the test data
#     y_hat_test2 = logreg_resampled.predict(X_test)

#     # Calculate accuracy
#     accuracy = accuracy_score(y_test, y_hat_test2)

#     # Calculate precision
#     precision = precision_score(y_test, y_hat_test2)

#     # Calculate recall
#     recall = recall_score(y_test, y_hat_test2)

#     # Calculate F1 score
#     f1 = f1_score(y_test, y_hat_test2)

#     # Create a confusion matrix
#     cm = confusion_matrix(y_test, y_hat_test2)
#     # Calculate the ROC curve
#     y_hat_prob2 = logreg_resampled.predict_proba(X_test)[: , 1]
#     fpr, tpr, thresholds = roc_curve(y_test, y_hat_prob2)

#     # Calculate the area under the ROC curve
#     auc = roc_auc_score(y_test, y_hat_prob2)

#     # Plot the ROC curve
#     plt.figure(figsize=(25,8))
#     plt.subplot(1,2,1)
#     plt.plot(fpr, tpr)
#     plt.xlabel("False Positive Rate")
#     plt.ylabel("True Positive Rate")
#     plt.title("ROC Curve")

#     # Create a heatmap of the confusion matrix
#     plt.subplot(1,2,2)
#     cm = confusion_matrix(y_test, y_hat_test2)
#     sns.heatmap(cm, annot=True, cmap='Blues', fmt='d')

#     # Add labels and title to the plot
#     plt.xlabel('Predicted')
#     plt.ylabel('Actual')
#     plt.title('Confusion Matrix')

#     # Display the plot
#     plt.show()
#     # Assuming the evaluation dictionary contains the evaluation metrics
#     evaluation = {
#         "Accuracy": accuracy,
#         "Precision": precision,
```

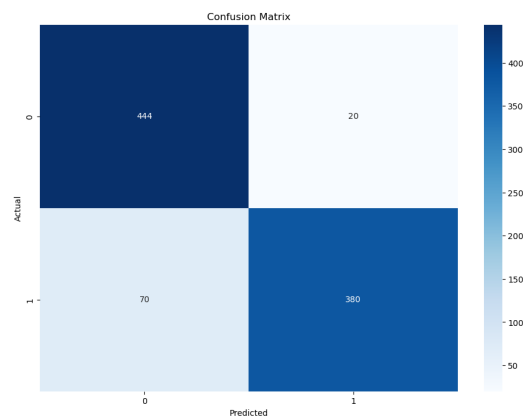
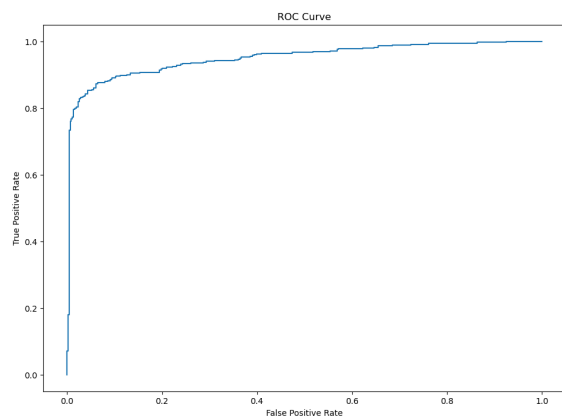
```
# "Recall": recall,
# "F1 Score": f1,
# "Confusion Matrix": cm,
# "AUC Score": auc
# }

# # Create a DataFrame from the evaluation dictionary
# evaluation_df = pd.DataFrame.from_dict(evaluation, orient='index', columns=['Metric'])

# # Display the DataFrame
# print(evaluation_df)
```

In [62]:

```
# evaluation2 = evaluate_logistic_regression(X_train_resampled, y_train_resampled)
# evaluation2
```



	Metrics
Accuracy	0.901532
Precision	0.95
Recall	0.844444
F1 Score	0.894118
Confusion Matrix	[[444, 20], [70, 380]]
AUC Score	0.950144

MODEL 2: K-NEAREST NEIGHBORS

The k-nearest neighbors (KNN) algorithm is a supervised machine learning method employed for classification and regression tasks. It estimates the probability of a data point belonging to a particular group by considering the group memberships of its nearest neighboring data points.

In [64]:

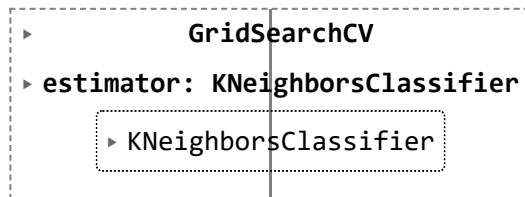


```

from sklearn.neighbors import KNeighborsClassifier
#instantiate the model
knn_clf= KNeighborsClassifier()
#set the paramerters for grid searchCV
knn_grid = {'n_neighbors':[1,2,3,4,5,6,7,8]}
#use GridSearchCV technique to search through the best parameter values
knn = GridSearchCV(knn_clf, knn_grid, scoring = 'accuracy', cv = 3 )
#fit the model
knn.fit(X_train, y_train.values.ravel())

```

Out[64]:



In [65]:



```

#use .best_params_ which tells use the best parameters for our model
knn.best_params_

```

Out[65]:

```
{'n_neighbors': 8}
```

In [66]:

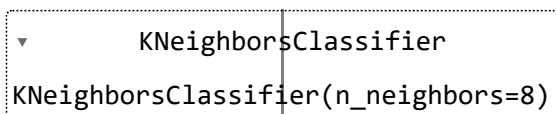


```

#Updating classifier with best parameters
knn = KNeighborsClassifier(n_neighbors = 8)
knn.fit(X,y.values.ravel())

```

Out[66]:



In [67]:



```
knn_prediction = knn.predict(X_test)
```

In [68]:



```
# Evaluating KNN Algorithm
print(confusion_matrix(y_test,knn_prediction))

TN, FP, FN, TP = confusion_matrix(y_test, knn_prediction).ravel()

print('True Positive(TP) = ', TP)
print('False Positive(FP) = ', FP)
print('True Negative(TN) = ', TN)
print('False Negative(FN) = ', FN)

accuracy = (TP+TN) / (TP+FP+TN+FN)

print('Accuracy of the binary classification = {:.3f}'.format(accuracy))
```

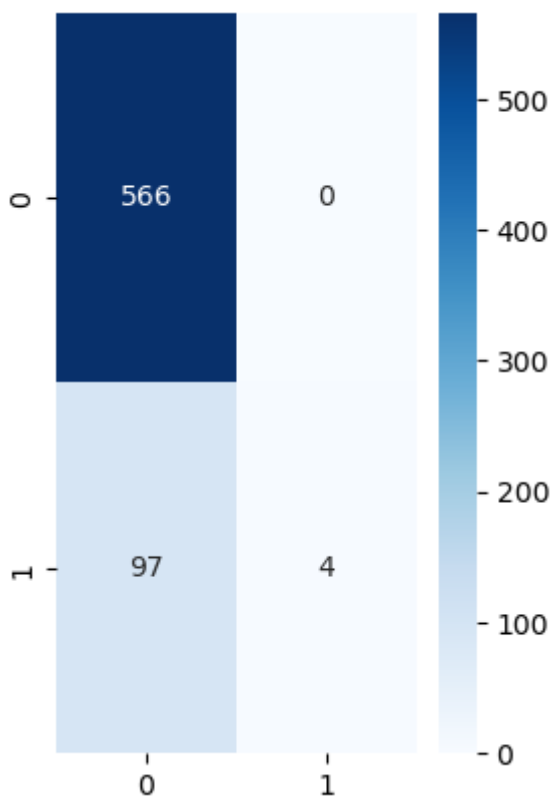
```
[[566   0]
 [ 97   4]]
True Positive(TP) = 4
False Positive(FP) = 0
True Negative(TN) = 566
False Negative(FN) = 97
Accuracy of the binary classification = 0.855
```

In [69]:

```
# Create a heatmap of the confusion matrix
plt.subplot(1,2,2)
cm = confusion_matrix(y_test,knn_prediction)
sns.heatmap(cm, annot=True, cmap='Blues', fmt='d')
```

Out[69]:

<Axes: >



Observations

Here's an explanation of the observations:

- True Positives (TP): The model correctly predicted the positive class for 4 instances. These are cases where the actual value is positive, and the model correctly identified them as positive.
- False Positives (FP): The model incorrectly predicted the positive class for 0 instances. These are cases where the actual value is negative, but the model mistakenly identified them as positive. In this case, there are no false positives.
- True Negatives (TN): The model correctly predicted the negative class for 566 instances. These are cases where the actual value is negative, and the model correctly identified them as negative.
- False Negatives (FN): The model incorrectly predicted the negative class for 97 instances. These are cases where the actual value is positive, but the model mistakenly identified them as negative.

The accuracy of the binary classification is calculated as the ratio of correct predictions (TP + TN) to the total number of instances. In this case, the accuracy is 0.855, which means that approximately 85.5% of instances were correctly classified by the model.

In [70]:



```
#model evaluation
knn_acc = accuracy_score(y_test, knn_prediction)
knn_f1 = f1_score(y_test, knn_prediction)
knn_prec = precision_score(y_test, knn_prediction)
knn_rec = recall_score(y_test, knn_prediction)
knn_auc = roc_auc_score(y_test, knn_prediction)

print('KNN:')
print(f'Accuracy: {knn_acc:.3f}')
print(f'F1 Score: {knn_f1:.3f}')
print(f'Precision: {knn_prec:.3f}')
print(f'Recall: {knn_rec:.3f}')
print(f'ROC AUC Score: {knn_auc:.3f}')
```

```
KNN:
Accuracy: 0.855
F1 Score: 0.076
Precision: 1.000
Recall: 0.040
ROC AUC Score: 0.520
```

Observations

Based on the provided observations of a K-Nearest Neighbors (KNN) model, here's an interpretation of the results:

1. Accuracy: The accuracy of 0.855 indicates that approximately 85.5% of the instances in the evaluation dataset were correctly classified by the KNN model. This metric alone, however, may not provide a comprehensive assessment of the model's performance.
2. F1 Score: The F1 score of 0.076 is a measure that balances both precision and recall. It indicates the harmonic mean of these two metrics. A low F1 score suggests poor performance in terms of correctly identifying positive instances and minimizing false positives.
3. Precision: The precision of 1.000 suggests that all instances predicted as positive by the KNN model were true positives. However, it is crucial to examine other metrics to assess the overall performance of the model.
4. Recall: The recall of 0.040 indicates that only a small proportion (approximately 4%) of actual positive instances were correctly identified by the KNN model. This implies that the model has a relatively high number of false negatives.
5. ROC AUC Score: The ROC AUC score of 0.520 represents the Area Under the Receiver Operating Characteristic Curve (ROC AUC). This metric assesses the model's ability to distinguish between positive and negative instances. An AUC score of 0.520 suggests that the model's discrimination capability is close to random guessing.

Overall, the observations indicate that the KNN model may have limitations in correctly identifying positive instances (low recall), and its overall performance in terms of precision, recall, and discrimination ability is relatively poor. Further analysis and model refinement may be necessary to improve its performance.

MODEL 3: DECISION TREES

Decision Trees (DTs) are a type of supervised learning technique used for classification and regression tasks. The objective is to build a model that can predict the value of a target variable based on simple decision rules learned from the features present in the data.

The term "non-parametric" indicates that decision trees do not rely on assumptions about the data's underlying distribution or specific parameter values. Instead, they focus on inferring straightforward decision rules directly from the data without imposing predefined assumptions about its characteristics or whether the data is quantitative or qualitative.

In [71]:



```
from sklearn.tree import DecisionTreeClassifier
#instantiate the model
tree = DecisionTreeClassifier(criterion="entropy", random_state=0)
#fit the data
tree.fit(X_train, y_train)
#predict
tree_y_hat = tree.predict(X_test)
#model evaluation using evaluation metrics
tree_acc = accuracy_score(y_test, tree_y_hat)
tree_f1 = f1_score(y_test, tree_y_hat)
tree_prec = precision_score(y_test, tree_y_hat)
tree_rec = recall_score(y_test, tree_y_hat)
tree_auc = roc_auc_score(y_test, tree_y_hat)
print('Decision Tree:')
print(f'Accuracy: {tree_acc:.3f}')
print(f'F1 Score: {tree_f1:.3f}')
print(f'Precision: {tree_prec:.3f}')
print(f'Recall: {tree_rec:.3f}')
print(f'ROC AUC Score: {tree_auc:.3f}')
```

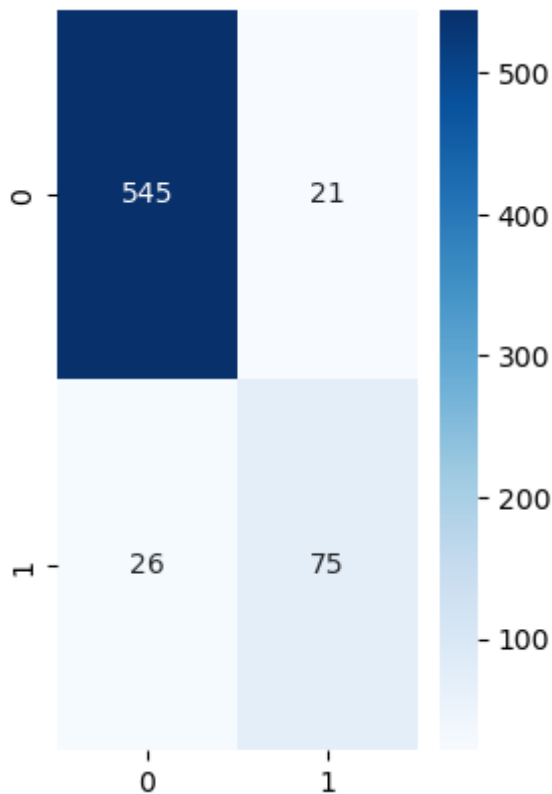
Decision Tree:
Accuracy: 0.930
F1 Score: 0.761
Precision: 0.781
Recall: 0.743
ROC AUC Score: 0.853

In [72]:

```
# Create a heatmap of the confusion matrix
plt.subplot(1,2,2)
cm = confusion_matrix(y_test,tree_y_hat)
sns.heatmap(cm, annot=True, cmap='Blues', fmt='d')
```

Out[72]:

<Axes: >



Observations

Based on the provided observations for a decision tree model, here's an interpretation of the results:

1. **Accuracy:** The accuracy of 0.93 indicates that approximately 93% of the instances in the evaluation dataset were correctly classified by the decision tree model. This is a relatively high accuracy, suggesting that the model performs well in terms of overall correct predictions.
2. **F1 Score:** The F1 score of 0.761 is a measure that balances both precision and recall. It represents the harmonic mean of these two metrics. A higher F1 score indicates a better balance between precision and recall. In this case, the F1 score suggests a moderate performance in correctly identifying positive instances and minimizing false positives.
3. **Precision:** The precision of 0.781 suggests that around 78.1% of the instances predicted as positive by the decision tree model are actually true positives. This metric measures the accuracy of positive predictions.
4. **Recall:** The recall of 0.743 indicates that approximately 74.3% of the actual positive instances were correctly identified by the decision tree model. This metric evaluates the model's ability to find all positive instances.
5. **ROC AUC Score:** The ROC AUC score of 0.853 represents the Area Under the Receiver Operating Characteristic Curve (ROC AUC). This metric assesses the model's ability to distinguish between positive and negative instances. A higher ROC AUC score suggests better discrimination capability. In

this case, the score of 0.853 indicates that the model has reasonably good discrimination ability.

Based on the provided confusion matrix for a decision tree model, here's an explanation of the observations:

- True Positives (TP): The model correctly predicted the positive class for 75 instances. These are cases where the actual value is positive, and the model correctly identified them as positive.
- False Positives (FP): The model incorrectly predicted the positive class for 21 instances. These are cases where the actual value is negative, but the model mistakenly identified them as positive.
- True Negatives (TN): The model correctly predicted the negative class for 545 instances. These are cases where the actual value is negative, and the model correctly identified them as negative.
- False Negatives (FN): The model incorrectly predicted the negative class for 26 instances. These are cases where the actual value is positive, but the model mistakenly identified them as negative.

Overall, the observations suggest that the decision tree model performs well in terms of accuracy, precision, recall, F1 score, and discrimination ability. However, further analysis and validation with additional evaluation metrics may be necessary to gain a more comprehensive understanding of the model's performance.

MODEL 4: RANDOM FOREST

In [73]:



```
from sklearn.ensemble import RandomForestClassifier
#instantiate the model
forest = RandomForestClassifier(n_estimators=10, criterion="entropy", random_state=0)
#fit the data
forest.fit(X_train, y_train)
#predict
forest_y_hat = forest.predict(X_test)
#Model evaluation using evaluation metrics
forest_acc = accuracy_score(y_test, forest_y_hat)
forest_f1 = f1_score(y_test, forest_y_hat)
forest_prec = precision_score(y_test, forest_y_hat)
forest_rec = recall_score(y_test, forest_y_hat)
forest_auc = roc_auc_score(y_test, forest_y_hat)
print('Random Forest:')
print(f'Accuracy: {forest_acc:.3f}')
print(f'F1 Score: {forest_f1:.3f}')
print(f'Precision: {forest_prec:.3f}')
print(f'Recall: {forest_rec:.3f}')
print(f'ROC AUC Score: {forest_auc:.3f}')
```

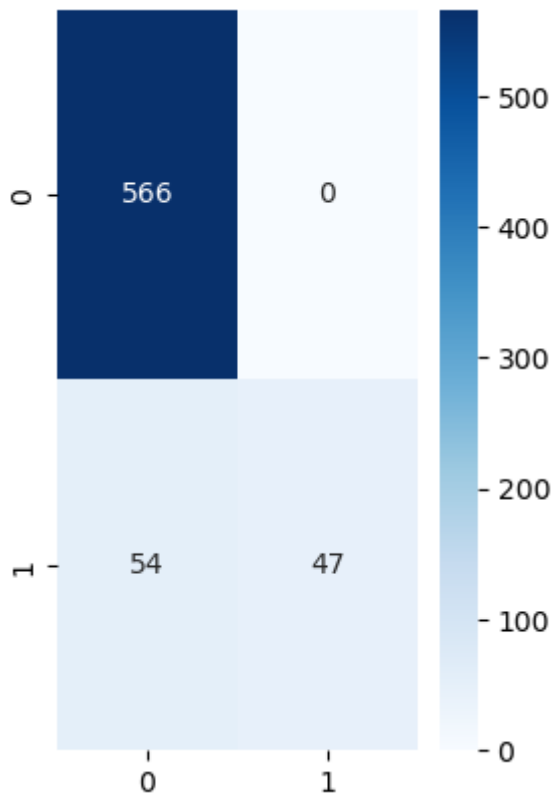
Random Forest:
Accuracy: 0.919
F1 Score: 0.635
Precision: 1.000
Recall: 0.465
ROC AUC Score: 0.733

In [74]:

```
# Create a heatmap of the confusion matrix
plt.subplot(1,2,2)
cm = confusion_matrix(y_test,forest_y_hat)
sns.heatmap(cm, annot=True, cmap='Blues', fmt='d')
```

Out[74]:

<Axes: >



Observations

Based on the provided observations for a random forest model, here's an interpretation of the results:

1. Accuracy: The accuracy of 0.919 indicates that approximately 91.9% of the instances in the evaluation dataset were correctly classified by the random forest model. This is a relatively high accuracy, suggesting that the model performs well in terms of overall correct predictions.
2. F1 Score: The F1 score of 0.635 is a measure that balances both precision and recall. It represents the harmonic mean of these two metrics. A higher F1 score indicates a better balance between precision and recall. In this case, the F1 score suggests a moderate performance in correctly identifying positive instances and minimizing false positives.
3. Precision: The precision of 1.00. This metric measures the accuracy of positive predictions.
4. Recall: The recall of 0.465 indicates that approximately 46.5% of the actual positive instances were correctly identified by the random forest model. This metric evaluates the model's ability to find all positive instances.
5. ROC AUC Score: The ROC AUC score of 0.733 represents the Area Under the Receiver Operating Characteristic Curve (ROC AUC). This metric assesses the model's ability to distinguish between positive and negative instances. A higher ROC AUC score suggests better discrimination capability. In this case, the score of 0.733 indicates that the model has some discrimination ability, but there is room for improvement.

Based on the provided confusion matrix for a random forest model, here's an explanation of the observations:

- True Positives (TP): The model correctly predicted the positive class for 47 instances. These are cases where the actual value is positive, and the model correctly identified them as positive.
- False Positives (FP): The model incorrectly predicted the positive class for 0 instances.
- True Negatives (TN): The model correctly predicted the negative class for 566 instances. These are cases where the actual value is negative, and the model correctly identified them as negative.
- False Negatives (FN): The model incorrectly predicted the negative class for 54 instances. These are cases where the actual value is positive, but the model mistakenly identified them as negative.

Overall, the observations suggest that the random forest model performs well in terms of accuracy and precision, indicating good overall predictions and accurate positive classifications. However, the model's performance in terms of recall and F1 score is relatively lower, suggesting room for improvement in correctly identifying positive instances and achieving a better balance between precision and recall.

MODEL SELECTION

In [75]:



```
# Compute the predicted probabilities for each model
knn_probability = knn.predict_proba(X_test)[:, 1]
log_probability = logreg.predict_proba(X_test)[:, 1]
tree_probability = tree.predict_proba(X_test)[:, 1]
forest_probability = forest.predict_proba(X_test)[:, 1]

# Compute the ROC curve and AUC for each model
tree_fpr, tree_tpr, _ = roc_curve(y_test, tree_probability)
tree_auc = roc_auc_score(y_test, tree_probability)

forest_fpr, forest_tpr, _ = roc_curve(y_test, forest_probability)
forest_auc = roc_auc_score(y_test, forest_probability)

knn_fpr, knn_tpr, _ = roc_curve(y_test, knn_probability)
knn_auc = roc_auc_score(y_test, knn_probability)

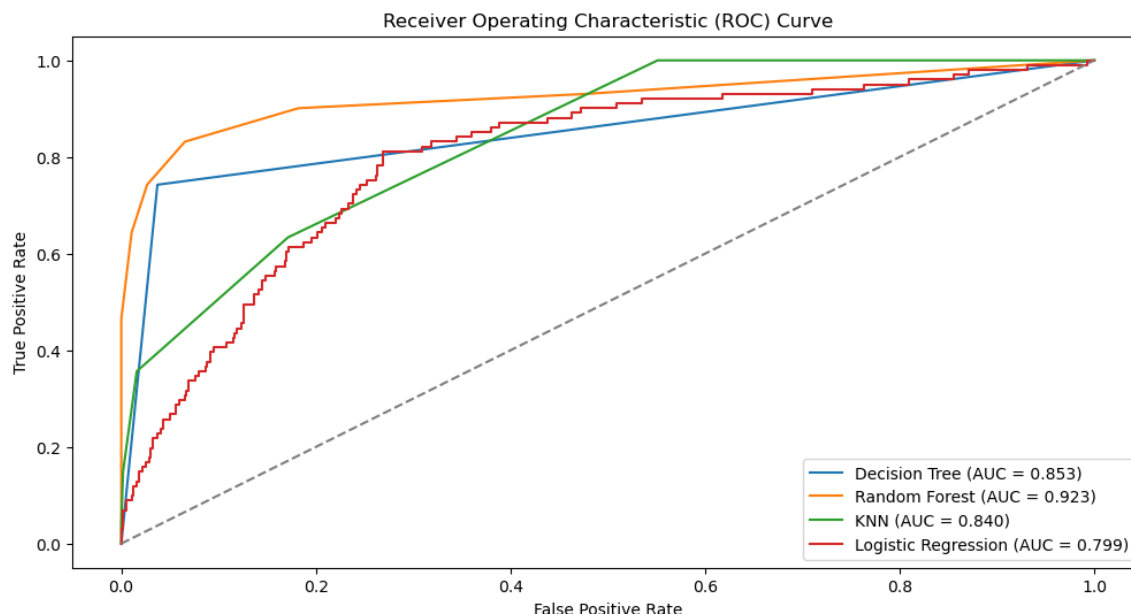
log_fpr, log_tpr, _ = roc_curve(y_test, log_probability)
log_auc = roc_auc_score(y_test, log_probability)

# Plot ROC curves
plt.figure(figsize=(12,6))
plt.plot(tree_fpr, tree_tpr, label='Decision Tree (AUC = {:.3f})'.format(tree_auc))
plt.plot(forest_fpr, forest_tpr, label='Random Forest (AUC = {:.3f})'.format(forest_auc))
plt.plot(knn_fpr, knn_tpr, label='KNN (AUC = {:.3f})'.format(knn_auc))
plt.plot(log_fpr, log_tpr, label='Logistic Regression (AUC = {:.3f})'.format(log_auc))
# Plot the random guess line
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')

# Set plot labels and title
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')

# Set plot legend
plt.legend()

# Show the plot
plt.show()
```



Observations

From the above evaluation metrics per model, we can draw the following conclusions:

- **Accuracy:** The Decision Tree and Random Forest models perform similarly well with accuracies of 0.930 and 0.919, respectively. Logistic Regression and KNN have slightly lower accuracies.
- **Precision:** Decision tree achieves the highest precision score of 0.781, indicating a high proportion of correct positive predictions. Logistic Regression has relatively lower precision score, while KNN and Random Forest achieves a perfect precision score of 1.000.
- **Recall:** The Decision Tree model has the highest recall score of 0.743, indicating its ability to identify a higher proportion of positive instances. Logistic Regression and Random Forest have relatively lower recall scores, while KNN performs the poorest in terms of recall.
- **F1 Score:** The Decision Tree model has the highest F1 score of 0.761, which considers both precision and recall. Random Forest follows. Logistic Regression and KNN have lower F1 scores, with KNN having the lowest.
- **ROC AUC Score:** The Decision Tree model achieves the highest ROC AUC score of 0.853, indicating its better ability to distinguish between positive and negative instances. Random Forest and Logistic Regression have relatively lower ROC AUC scores, while KNN has the lowest score.

In summary, the Decision Tree and Random Forest models generally perform better across multiple evaluation metrics, including accuracy, precision, recall, F1 score, and ROC AUC score. Logistic Regression performs moderately, while KNN shows relatively lower performance in most of the evaluation metrics.

FEATURE IMPORTANCE

In [76]:

```
# Define the list of columns to drop
columns_to_drop = ['state_AL', 'state_AR', 'state_AZ', 'state_CA', 'state_CO', 'state_CT',
                   'state_DE', 'state_FL', 'state_GA', 'state_HI', 'state_IA', 'state_ID',
                   'state_IN', 'state_KS', 'state_KY', 'state_LA', 'state_MA', 'state_MD',
                   'state_MI', 'state_MN', 'state_MO', 'state_MS', 'state_MT', 'state_NC',
                   'state_NE', 'state_NH', 'state_NJ', 'state_NM', 'state_NV', 'state_NY',
                   'state_OK', 'state_OR', 'state_PA', 'state_RI', 'state_SC', 'state_SD',
                   'state_TX', 'state_UT', 'state_VA', 'state_VT', 'state_WA', 'state_WI',
                   'state_WY']

# Drop the specified columns from the DataFrame
df2 = df2.drop(columns=columns_to_drop)
```

In [77]:

```
# Split the DataFrame into input features 'X' and target variable 'y'
y = df2['churn']
X = df2.drop('churn', axis = 1)

# Train a decision tree classifier
tree = DecisionTreeClassifier()
tree.fit(X, y)

# Get feature importances
importance = tree.feature_importances_

# Print feature importances
for i, feature in enumerate(df2.columns[:-1]): # Exclude the target variable column
    print(f"{feature}: {importance[i]}")
```

```
account length: 0.035491157983519976
area code: 0.01420242895022235
number vmail messages: 0.048659376641064574
total day calls: 0.02580171736752936
total day charge: 0.2711185176412581
total eve calls: 0.02173415017450898
total eve charge: 0.14074433916148968
total night calls: 0.020086035174988578
total night charge: 0.046642178522023875
total intl calls: 0.06396022434240756
total intl charge: 0.10548427399782369
customer service calls: 0.11151218376317397
churn: 0.07235405876074377
international plan_yes: 0.022209357519245514
```

In [79]:

```

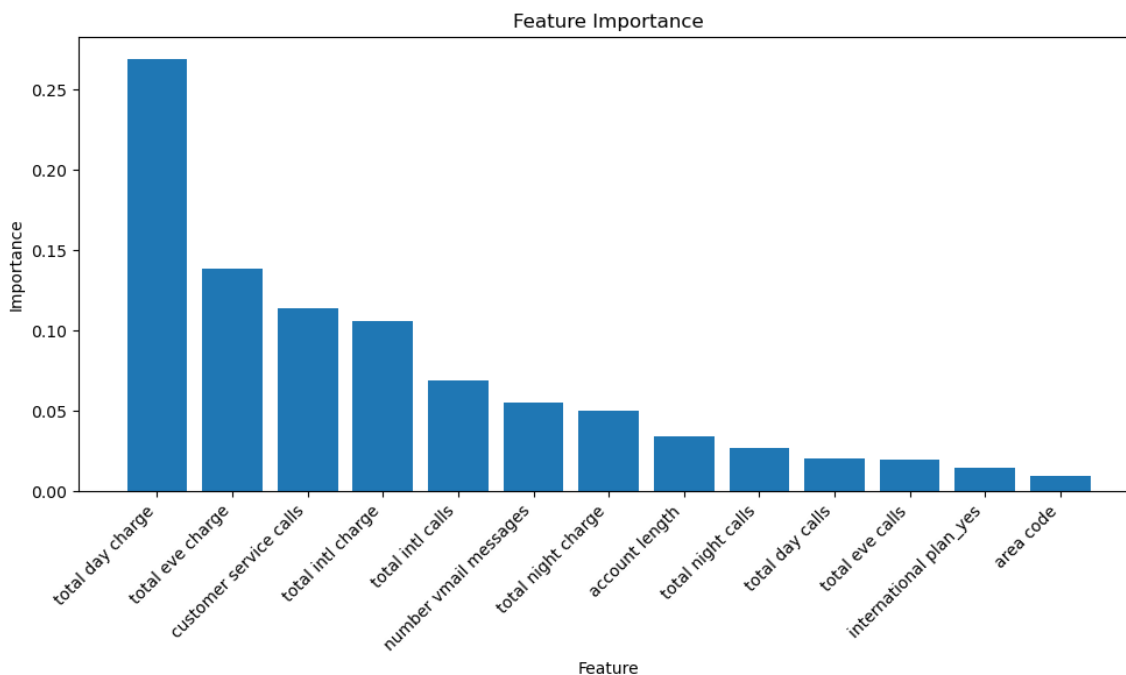
import matplotlib.pyplot as plt
import numpy as np

# Define the feature names and importances
feature_names = ['account length', 'area code', 'number vmail messages', 'total day calls',
                 'total eve calls', 'total eve charge', 'total night calls', 'total night charge',
                 'total intl calls', 'total intl charge', 'customer service calls', 'international plan_yes',
                 'area code']
importances = [0.034251971990780006, 0.009754261429972448, 0.055155819277019416, 0.020285117641845,
               0.26898307117641845, 0.019562760427114006, 0.13841738365359268, 0.02678195117641845,
               0.050274084373609355, 0.06916595606301343, 0.10618930716386868, 0.11373168117641845,
               0.009754261429972448]

# Sort the features and importances in descending order
sorted_indices = np.argsort(importances)[::-1]
sorted_feature_names = [feature_names[i] for i in sorted_indices]
sorted_importances = np.sort(importances)[::-1]

# Plot the feature importances
plt.figure(figsize=(10, 6))
plt.bar(range(len(sorted_importances)), sorted_importances, align='center')
plt.xticks(range(len(sorted_importances)), sorted_feature_names, rotation=45, ha='right')
plt.xlabel('Feature')
plt.ylabel('Importance')
plt.title('Feature Importance')
plt.tight_layout()
plt.show()

```



- The most important feature for predicting churn is **total day charge**, which has a score of 0.268983. This means that the amount of money a customer spends on day calls is a strong predictor of whether they will churn.
- The second most important feature is **customer service calls**, which has a score of 0.113732. This means that customers who make more customer service calls are more likely to churn.
- Other important features include **total eve charge**, **total intl charge**, and **international plan_yes**. These features all relate to the amount of money a customer spends on their phone service, which is a strong predictor of churn.

- The least important features are **account length**, **area code**, and **number vmail messages**. These features do not seem to be very predictive of churn.

Overall, the feature importance indicates that the amount of money a customer spends on their phone service is a strong predictor of whether they will churn. Other important factors include the number of customer service calls a customer makes and whether they have an international plan.

CONCLUSION

Decision Tree model appears to be the best performer among the four models. This would be the best Model for the Syria Tel Telecommunication Company to use to predict which customer will unsubscribe from their services and take precautionary steps to reduce the churn rate.

The Most important features for predicting churn are:

- Total day charge
- Customer Service call
- Total eve charge
- Total intl charge

RECOMMENDATIONS

- **Focus on reducing the amount of money customers spend on day calls.** This is the most important factor in predicting churn, so it is the most important area to focus on. This could be done by offering discounts on day calls, or by providing customers with more affordable alternatives to day calls.
- **Reduce the number of customer service calls.** Customers who make more customer service calls are more likely to churn. This could be done by improving the customer service experience, or by making it easier for customers to resolve their issues without having to call customer service.
- **Consider offering international plans.** Customers who have international plans are less likely to churn. This could be done by offering more affordable international plans, or by making it easier for customers to sign up for international plans.
- **Ignore account length, area code, and number vmail messages.** These features are not very predictive of churn, so there is no need to focus on them.

In []:

